

Cmput 291

Project 2

Sarah Morris, Victoria Bobey, Eldon Lake

Experimental Results

Hash Table times:

Search #/Search Type	Key	Data	Range
1	5.1021575927734375	1086629042.1485901	111.38916015625
2	6.413459777832031	366598895.072937	
3	4.792213439941406	1265428450.1075745	
4	5.412101745605469	1431917526.960373	
5	4.9114227294921875	311798863.17253113	
Average	5.326274843	836350755.5	

B-Tree Times:

Search #/Search Type	Key	Data	Range
1	7.700920104980469	539477543.592453	116.20521545410156
2	4.696846008300781	964701429.8439026	97.70393371582031
3	6.890296936035156	461164918.422699	90.28911590576172
4	7.2002410888671875	539477543.592453	121.80805206298828
5	7.0095062255859375	862524826.7650604	102.71072387695312
Average	6.699562073	673469252.4	105.9447008

Index File times:

Search # /Search Type	Key	Data	Range
1	45.490264892578125	30.58910369873047	82.89813995361328
2	5.984306335449219	36.88335418701172	121.9034194946289
3	5.698204040527344	35.71510314941406	89.00165557861328
4	5.793571472167969	31.614303588867188	79.29801940917969
5	5.602836608886719	5.412101745605469	104.28428649902344
Average	13.71383549	15.39707177	76.70593261

IndexFile:

Our IndexFile consists of a normal b-tree table that is used for key search and range search and a secondary index b-tree that is used for data search.

The b-tree was fast with both key and range search but very slow with the data search. The hash table proved to be no better at this. The secondary index allows the data search to be preformed in much the

same way that the key search is preformed, but instead of the keys being the index that is being searched through, the data is.

Analysis:

The hash table proved to be efficient for the key search. On average it only took 5 microseconds. The data search was much slower, averaging a runtime of 83635075.55 microseconds. The range search is also quit slow. It took on average **(INSERT AVERAGE HERE)** microseconds.

The b-tree proved to be efficient for both key and range search but very slow for data search. Key search can be preformed with an average time of 5 microseconds. Range search can be performed with an average time of 106 microseconds. The data search took longer at an average time of 673469252.4 microseconds.

The data search should be much slower for the hash table and the b-tree since each key/data combination must be checked until one is found that matches the given data. This could potentially mean that all of the key/data pairs must be checked.

The range search was slow for the hash table since the data is stored in no order. The program must search through all of the keys looking for any that fall between the given upper and lower bounds. The b-tree did not have to do this since the data is stored in more of a sequential order.

Since the b-tree was able to preform range searches significantly faster than the hash table, we choose to make our indexFile a b-tree. The b-tree preforms fast key searches as well as range searches. The data search was preformed on a b-tree with a secondary index. Having this allows for the completion of all three searches in very little time.

The secondary index in the indexFile allows the program to run through the b-tree like any other key search but instead of the keys being the normal keys from the key/data pair the key is the data portion of this pair. This means that the data search is just as fast as the key search. This is a significant improvement from the hash table and b-tree (with no secondary index) performance of the data search.

As you can see from the experimental results the key search and data search that were preformed by the indexFile have very similar averages for time. Both are very quick with run time averages of 14 and 15 microseconds respectively. The Range search takes more time, but is still very fast at only 77 microseconds on average.